

# Laboratorio 4: Clasificación de Imágenes con CNN (CIFAR-10)

## Integrantes

- José Rodrigo Marchena, 22398
- Sofía Velasquez, 22049

## Objetivo general

Implementar y comparar un modelo base ANN y una CNN para clasificar el dataset CIFAR-10, evaluando desempeño, curvas de aprendizaje y errores comunes. También se incluye un ejercicio opcional de Data Augmentation.

---

## Parte 1: Preparación del Conjunto de Datos

1. Cargamos el dataset CIFAR-10 directamente desde `keras.datasets`.
  - Son 60,000 imágenes de tamaño 32x32x3 (RGB).
  - Se dividen en 50,000 de entrenamiento y 10,000 de prueba.
2. Normalizamos los valores de píxeles, dividiéndolos entre 255, de modo que los valores queden en el rango [0,1]. Cómo las imágenes son arrays entre 0 y 255 de la intensidad de cada pixel.
3. Definimos los nombres de las clases (avión, coche, perro, etc.).
4. Mostramos 10 imágenes de entrenamiento con sus etiquetas para confirmar que los datos están bien cargados.

```
In [1]: from matplotlib import pyplot as plt
import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.datasets import cifar10

(x_train, y_train), (x_test, y_test) = cifar10.load_data() # Carga de CIFAR-10 data

class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck'] # Definir los nombres de las cl

# Mostrar las formas de los conjuntos de datos
print("Shapes -> x_train:", x_train.shape, "y_train:", y_train.shape)
print("Shapes -> x_test:", x_test.shape, "y_test:", y_test.shape)

# Normalización [0,1]
x_train = x_train.astype("float32") / 255.0
x_test = x_test.astype("float32") / 255.0
```

```
# Visualización de 10 ejemplos
def show_grid(images, labels, n=10):
    plt.figure(figsize=(8,8))
    idx = np.random.choice(len(images), n, replace=False)
    for i, j in enumerate(idx):
        plt.subplot(5,5,i+1)
        plt.imshow(images[j])
        plt.title(class_names[labels[j][0]])
        plt.axis("off")
    plt.tight_layout()
    plt.show()

show_grid(x_train, y_train)
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
**170498071/170498071** ————— **44s** 0us/step  
 Shapes -> x\_train: (50000, 32, 32, 3) y\_train: (50000, 1)  
 Shapes -> x\_test: (10000, 32, 32, 3) y\_test: (10000, 1)



## Parte 2: Modelo Base ANN

Se define una función `build_ann` que implementa una red neuronal artificial densa (ANN):

- Se aplanan la imagen ( $32 \times 32 \times 3 \rightarrow 3072$  neuronas).
- Dos capas densas ocultas con activación ReLU.
- Capa final Dense(10) con activación Softmax para clasificar las 10 clases.

Se compila el modelo con:

- **Optimizer:** Adam.
- **Loss:** categorical\_crossentropy.
- **Métrica:** Accuracy.

Luego lo entrenamos, medimos tiempo de entrenamiento y evaluamos el rendimiento en el set de prueba.

```
In [2]: import time
from tensorflow.keras.utils import to_categorical

# One-hot encoding
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

# Función para construir el modelo ANN
def build_ann(input_shape=(32,32,3), num_classes=10):
    model = keras.Sequential([
        layers.Flatten(input_shape=input_shape), # Aplanar la imagen 32x32x3 a un v
        layers.Dense(512, activation="relu"), # Capa oculta con 512 neuronas y acti
        layers.Dense(256, activation="relu"), # Capa oculta con 256 neuronas y acti
        layers.Dense(num_classes, activation="softmax") # Capa de salida con activa
    ])
    model.compile(optimizer='adam',
                  loss='categorical_crossentropy', # Función de pérdida para clasificac
                  metrics=['accuracy']) # Métrica de precisión
    return model

# Construcción del modelo ANN
ann = build_ann()
# Mostrar resumen del modelo
ann.summary()
```

C:\Users\JM\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11\_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\keras\src\layers\reshaping\flatten.py:37: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(**kwargs)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 3072)	0
dense (Dense)	(None, 512)	1,573,376
dense_1 (Dense)	(None, 256)	131,328
dense_2 (Dense)	(None, 10)	2,570

Total params: 1,707,274 (6.51 MB)

Trainable params: 1,707,274 (6.51 MB)

Non-trainable params: 0 (0.00 B)

## Entrenamiento del modelo ANN

En este bloque de código se entrena el modelo(ANN): Se define:

- `batch_size = 64` que va a ser el tamaño del lote

- `epochs = 10` que van a ser las épocas que va a recorrer el modelo

Para el entrenamiento se le pasan los parámetros a la función para entrenar el modelo:

- Se le mandan los conjuntos de entrenamiento
- El `validation_split=0.2` para usar 20% de los datos de entrenamiento para validación
- El `batch_size` y `epochs`
- El `verbose=2` para mostrar información detallada durante el entrenamiento

```
In [ ]: # Entrenamiento ANN
ann = build_ann()
batch_size = 64
epochs = 10

start = time.time() # Inicio del temporizador
hist_ann = ann.fit(
    x_train, y_train,
    validation_split=0.2, # Usar 20% de los datos de entrenamiento para validación
    batch_size=batch_size,
    epochs=epochs,
    verbose=2 # Mostrar información detallada durante el entrenamiento
)
t_train_ann = time.time() - start # fin del temporizador
# Mostrar tiempo de entrenamiento
print(f"Tiempo de entrenamiento ANN: {t_train_ann:.2f} s")

# Evaluación del modelo en test
test_loss_ann, test_acc_ann = ann.evaluate(x_test, y_test, verbose=0)
print(f"[ANN] Test Acc: {test_acc_ann:.4f} | Test Loss: {test_loss_ann:.4f}")
```

```

Epoch 1/10
625/625 - 18s - 28ms/step - accuracy: 0.3171 - loss: 1.9101 - val_accuracy: 0.3634 -
val_loss: 1.7818
Epoch 2/10
625/625 - 16s - 26ms/step - accuracy: 0.3871 - loss: 1.7027 - val_accuracy: 0.4085 -
val_loss: 1.6564
Epoch 3/10
625/625 - 16s - 26ms/step - accuracy: 0.4251 - loss: 1.6105 - val_accuracy: 0.4203 -
val_loss: 1.6307
Epoch 4/10
625/625 - 16s - 26ms/step - accuracy: 0.4449 - loss: 1.5484 - val_accuracy: 0.4392 -
val_loss: 1.5872
Epoch 5/10
625/625 - 16s - 25ms/step - accuracy: 0.4602 - loss: 1.5079 - val_accuracy: 0.4412 -
val_loss: 1.5749
Epoch 6/10
625/625 - 16s - 25ms/step - accuracy: 0.4768 - loss: 1.4652 - val_accuracy: 0.4671 -
val_loss: 1.5175
Epoch 7/10
625/625 - 16s - 25ms/step - accuracy: 0.4835 - loss: 1.4448 - val_accuracy: 0.4629 -
val_loss: 1.5391
Epoch 8/10
625/625 - 16s - 25ms/step - accuracy: 0.4943 - loss: 1.4155 - val_accuracy: 0.4673 -
val_loss: 1.5122
Epoch 9/10
625/625 - 16s - 25ms/step - accuracy: 0.5017 - loss: 1.3949 - val_accuracy: 0.4751 -
val_loss: 1.5004
Epoch 10/10
625/625 - 16s - 25ms/step - accuracy: 0.5098 - loss: 1.3662 - val_accuracy: 0.4787 -
val_loss: 1.4690
Tiempo de entrenamiento ANN: 161.29 s
[ANN] Test Acc: 0.4809 | Test Loss: 1.4493

```

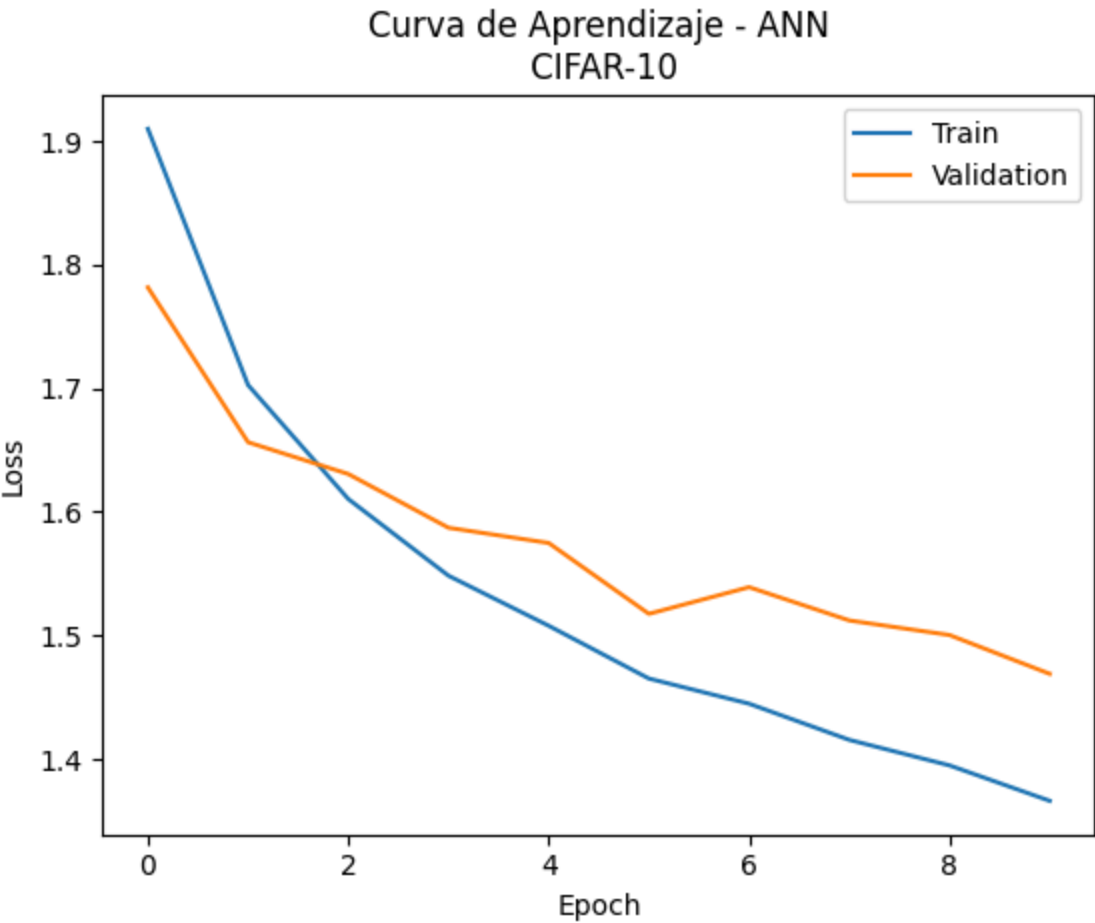
```

In [47]: from matplotlib import pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
# Learning Curve
plt.title("Curva de Aprendizaje - ANN\n CIFAR-10")
plt.plot(np.arange(0,epochs), hist_ann.history["loss"], label="Train")
plt.plot(np.arange(0,epochs), hist_ann.history["val_loss"], label="Validation")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.show()

# Confusion Matrix
pred_start = time.time()
y_pred_prob = ann.predict(x_test)
print(f"Tiempo de prediccion {time.time() -pred_start}")

y_pred_class = np.argmax(y_pred_prob,axis=1)
y_test_class = np.argmax(y_test, axis=1)
print(classification_report(y_test_class, y_pred_class))
cm = confusion_matrix(y_test_class, y_pred_class)
sns.heatmap(cm, xticklabels=class_names, yticklabels=class_names)
plt.title("ANN - Matriz de Confusion\n CIFAR-10")

```

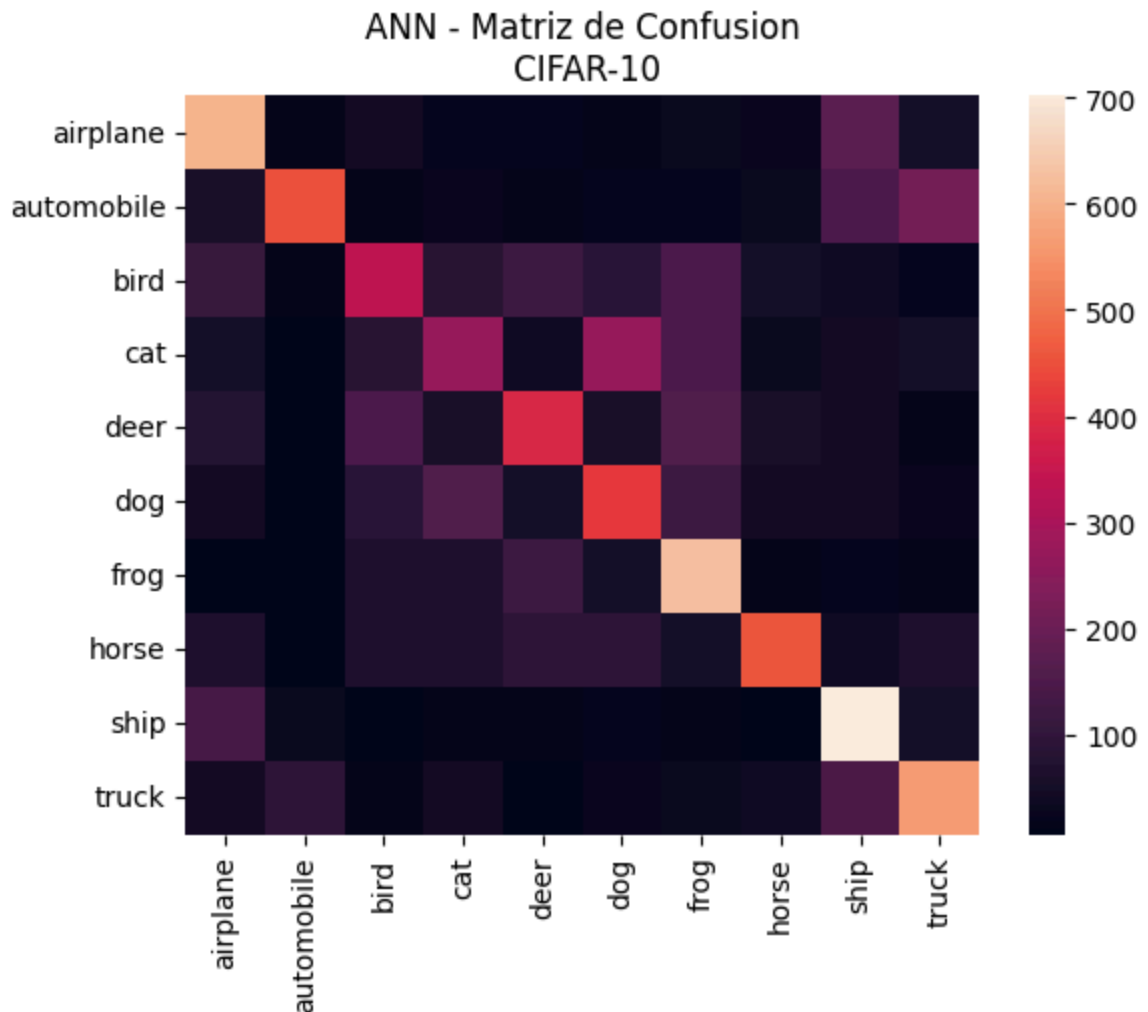


313/313 ————— 1s 3ms/step

Tiempo de prediccion 1.0924034118652344

	precision	recall	f1-score	support
0	0.50	0.60	0.55	1000
1	0.70	0.45	0.55	1000
2	0.39	0.33	0.36	1000
3	0.33	0.27	0.30	1000
4	0.45	0.38	0.41	1000
5	0.39	0.42	0.40	1000
6	0.46	0.63	0.53	1000
7	0.62	0.46	0.53	1000
8	0.50	0.70	0.59	1000
9	0.53	0.57	0.55	1000
accuracy			0.48	10000
macro avg	0.49	0.48	0.48	10000
weighted avg	0.49	0.48	0.48	10000

Out[47]: Text(0.5, 1.0, 'ANN - Matriz de Confusion\n CIFAR-10')



## Parte 3: Implementación de CNN

Se define una función `build_cnn` para diseñar una red neuronal convolucional (CNN):

- Varias capas Conv2D con filtros (32, 64, 128).
- Capas de MaxPooling2D para reducir la dimensión espacial.
- Dropout en puntos estratégicos para reducir sobreajuste.
- Al final, una capa densa y salida Softmax de 10 clases.

Luego se registra el tiempo de entrenamiento y se evalúa con los datos de prueba.

```
In [43]: from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping

# Función para construir el modelo CNN
def build_cnn(input_shape=(32,32,3), num_classes=10, dropout_rate=0.3):
    model = keras.Sequential([
        Conv2D(32, (3,3), activation='relu', padding='same', input_shape=input_shape),
        MaxPooling2D((2,2)), # Capa de max pooling 1
        Dropout(dropout_rate),
        Conv2D(64, (3,3), activation='relu', padding='same'), # Capa convolucional
```

```

        MaxPooling2D((2,2)), # Capa de max pooling 2
        Dropout(dropout_rate),
        Conv2D(128, (3,3), activation='relu', padding='same'), # Capa convolucional
        MaxPooling2D((2,2)), # Capa de max pooling 3
        Dropout(dropout_rate),
        Flatten(), # Aplanar las características extraídas
        Dense(256, activation='relu'), # Capa densa oculta con 256 neuronas y activ
        Dropout(dropout_rate),
        Dense(num_classes, activation='softmax') # Capa de salida con activación s
    ])
    model.compile( # Compilación del modelo
        optimizer='adam',
        loss="categorical_crossentropy", # Función de pérdida para clasificación mu
        metrics=["accuracy"] # Métrica de precisión
    )
    return model

# Construcción del modelo CNN
cnn = build_cnn()
# Mostrar resumen del modelo
cnn.summary()

```

C:\Users\JM\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11\_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\keras\src\layers\convolutional\base\_conv.py:113: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential\_1"



Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
dropout (Dropout)	(None, 16, 16, 32)	0
conv2d_1 (Conv2D)	(None, 16, 16, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_1 (Dropout)	(None, 8, 8, 64)	0
conv2d_2 (Conv2D)	(None, 8, 8, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
dropout_2 (Dropout)	(None, 4, 4, 128)	0
flatten_1 (Flatten)	(None, 2048)	0
dense_3 (Dense)	(None, 256)	524,544
dropout_3 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 10)	2,570

Total params: 620,362 (2.37 MB)

Trainable params: 620,362 (2.37 MB)

Non-trainable params: 0 (0.00 B)

## Entrenamiento del modelo CNN

En este bloque de código se entrena el modelo(ANN):

Se llama a la función para entrenar el modelo:

- Se mandan los conjuntos de entrenamiento, `epochs = 10` , `validation_split=0.2` , `batch_size = 64` , `verbose=2` que es exactamente lo mismo que se le mandó al modelo ann
- Solo se le agrega `callbacks=[early]`
  - Entrenamos el modelo usando validación y `EarlyStopping` para detener si no mejora la accuracy en validación.

```
In [44]: early = keras.callbacks.EarlyStopping( # Callback para detener el entrenamiento tem
monitor="val_accuracy", # Monitorear la precisión en el conjunto de validación
patience=5, # Paciencia para detener el entrenamiento
restore_best_weights=True # Restaurar Los mejores pesos al detener el entrenami
)
```

```

start = time.time() # Inicio del temporizador
# Entrenamiento del modelo CNN
hist_cnn = cnn.fit(
    x_train, y_train, # Datos de entrenamiento
    validation_split=0.2, # Usar 20% de los datos de entrenamiento para validación
    batch_size=64, # Tamaño del lote
    epochs=10, #
    callbacks=[early],
    verbose=2
)
t_train_cnn = time.time() - start # Tiempo de entrenamiento
print(f"Tiempo de entrenamiento CNN: {t_train_cnn:.2f} s")

test_loss_cnn, test_acc_cnn = cnn.evaluate(x_test, y_test, verbose=0)
print(f"[CNN] Test Acc: {test_acc_cnn:.4f} | Test Loss: {test_loss_cnn:.4f}")

```

Epoch 1/10

625/625 - 28s - 45ms/step - accuracy: 0.3816 - loss: 1.6806 - val\_accuracy: 0.5168 - val\_loss: 1.3471

Epoch 2/10

625/625 - 26s - 41ms/step - accuracy: 0.5312 - loss: 1.3001 - val\_accuracy: 0.6029 - val\_loss: 1.1240

Epoch 3/10

625/625 - 29s - 47ms/step - accuracy: 0.5871 - loss: 1.1570 - val\_accuracy: 0.6512 - val\_loss: 0.9917

Epoch 4/10

625/625 - 29s - 47ms/step - accuracy: 0.6246 - loss: 1.0568 - val\_accuracy: 0.6769 - val\_loss: 0.9250

Epoch 5/10

625/625 - 28s - 45ms/step - accuracy: 0.6486 - loss: 0.9901 - val\_accuracy: 0.6976 - val\_loss: 0.8747

Epoch 6/10

625/625 - 29s - 47ms/step - accuracy: 0.6687 - loss: 0.9328 - val\_accuracy: 0.7092 - val\_loss: 0.8355

Epoch 7/10

625/625 - 28s - 45ms/step - accuracy: 0.6877 - loss: 0.8935 - val\_accuracy: 0.7095 - val\_loss: 0.8203

Epoch 8/10

625/625 - 29s - 46ms/step - accuracy: 0.6971 - loss: 0.8539 - val\_accuracy: 0.7348 - val\_loss: 0.7713

Epoch 9/10

625/625 - 29s - 47ms/step - accuracy: 0.7085 - loss: 0.8241 - val\_accuracy: 0.7346 - val\_loss: 0.7703

Epoch 10/10

625/625 - 29s - 47ms/step - accuracy: 0.7199 - loss: 0.7898 - val\_accuracy: 0.7492 - val\_loss: 0.7227

Tiempo de entrenamiento CNN: 286.61 s

[CNN] Test Acc: 0.7476 | Test Loss: 0.7378

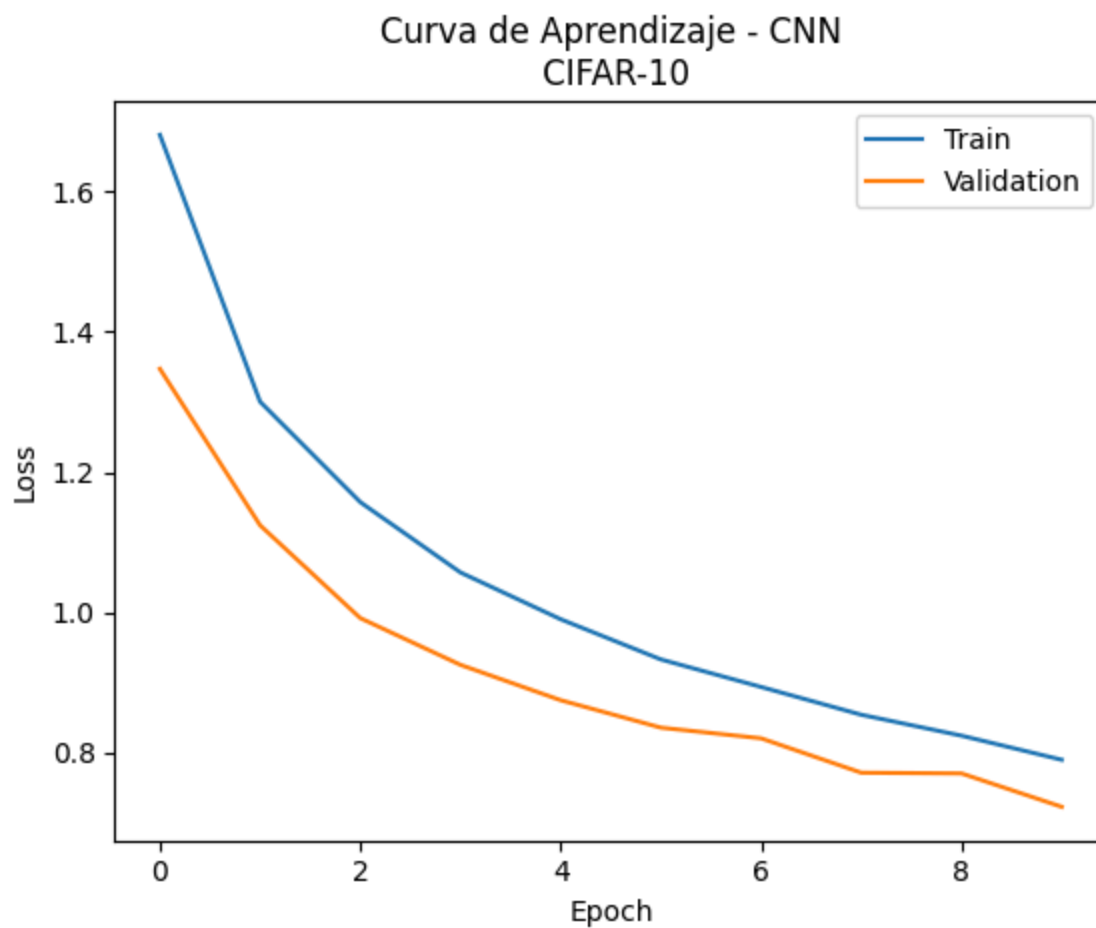
```

In [48]: # Learning curve
plt.title("Curva de Aprendizaje - CNN\n CIFAR-10")
plt.plot(np.arange(0, epochs), hist_cnn.history["loss"], label="Train")
plt.plot(np.arange(0, epochs), hist_cnn.history["val_loss"], label="Validation")
plt.xlabel("Epoch")
plt.ylabel("Loss")

```

```
plt.legend()
plt.show()

# Confusion Matrix
pred_start = time.time()
y_pred_prob = cnn.predict(x_test)
print(f"Tiempo de prediccion {time.time() -pred_start}")
y_pred_class = np.argmax(y_pred_prob,axis=1)
y_test_class = np.argmax(y_test, axis=1)
print(classification_report(y_test_class, y_pred_class))
cm = confusion_matrix(y_test_class, y_pred_class)
sns.heatmap(cm, xticklabels=class_names, yticklabels=class_names)
plt.title("CNN - Matriz de Confusion\n CIFAR-10")
```




313/313 ————— 2s 6ms/step

Tiempo de prediccion 2.1732585430145264

	precision	recall	f1-score	support
0	0.76	0.81	0.78	1000
1	0.86	0.88	0.87	1000
2	0.62	0.66	0.64	1000
3	0.63	0.46	0.53	1000
4	0.69	0.68	0.69	1000
5	0.62	0.70	0.66	1000
6	0.75	0.87	0.81	1000
7	0.84	0.76	0.80	1000
8	0.85	0.86	0.86	1000
9	0.85	0.81	0.83	1000
accuracy			0.75	10000
macro avg	0.75	0.75	0.75	10000
weighted avg	0.75	0.75	0.75	10000

Out[48]: Text(0.5, 1.0, 'CNN - Matriz de Confusion\n CIFAR-10')



## Parte 4: Ejemplificacion de errores

```
In [94]: def show_prediction_errors(model, model_name, wrong_count):
# Init
wrong_images = []
wrong_labels = []
correct_labels = []
counter = 0

# Find error examples
while(len(wrong_labels) < wrong_count):
    x_single = np.expand_dims(x_test[counter], axis=0)
    pred_prob = model.predict(x_single, verbose=0)
    pred_class = np.argmax(pred_prob)

    real_class = np.argmax(y_test[counter])

    if(pred_class != real_class):
        wrong_images.append(np.squeeze(x_single, axis=0) )
        wrong_labels.append(pred_class)
        correct_labels.append(real_class)
        counter+=1

# Show error examples
plt.figure(figsize=(8,8))
plt.suptitle(f"Error Examples - {model_name}")
for i in range(wrong_count):
    plt.subplot(5,5,i+1)
    plt.imshow(wrong_images[i])
    plt.title(
        f"Real: {class_names[correct_labels[i]]}\nPredicted: {class_names[wrong_labels[i]]}"
    )
    plt.axis("off")
plt.tight_layout()
plt.show()
```

```
In [ ]: # Para ANN
show_prediction_errors(ann, "ANN", 5)
```

Error Examples - ANN



```
In [ ]: # Para CNN
show_prediction_errors(cnn, "CNN", 5)
```

## Error Examples - CNN



## Puntos Extra: Data Augmentation

```
In [107... from tensorflow.keras.layers import RandomFlip, RandomRotation, RandomZoom

def build_cnn_augmented(input_shape=(32,32,3), num_classes=10, dropout_rate=0.3):
    model = keras.Sequential([
        # Data augmentation
        RandomFlip("horizontal_and_vertical", input_shape=input_shape),
        RandomZoom(0.1),
        RandomRotation(0.1),

        # Convolutional Layers
        Conv2D(32, (3,3), activation='relu', padding='same', input_shape=input_shape),
        MaxPooling2D((2,2)), # Capa de max pooling 1
        Dropout(dropout_rate),
        Conv2D(64, (3,3), activation='relu', padding='same'), # Capa convolucional
        MaxPooling2D((2,2)), # Capa de max pooling 2
        Dropout(dropout_rate),
        Conv2D(128, (3,3), activation='relu', padding='same'), # Capa convolucional
        MaxPooling2D((2,2)), # Capa de max pooling 3
        Dropout(dropout_rate),

        # Fully Connected Layers
        Flatten(), # Aplanar las características extraídas
        Dense(256, activation='relu'), # Capa densa oculta con 256 neuronas y activación
        Dropout(dropout_rate),
        Dense(num_classes, activation='softmax') # Capa de salida con activación s
    ])
    model.compile( # Compilación del modelo
        optimizer='adam',
        loss="categorical_crossentropy", # Función de pérdida para clasificación mu
        metrics=["accuracy"] # Métrica de precisión
    )
    return model

# Construcción del modelo CNN
cnn_aug = build_cnn_augmented()
# Mostrar resumen del modelo
cnn_aug.summary()
```

C:\Users\JM\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11\_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\keras\src\layers\preprocessing\tf\_data\_layer.py:19: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(**kwargs)
```

C:\Users\JM\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11\_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\keras\src\layers\convolutional\base\_conv.py:113: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential\_12"

Layer (type)	Output Shape	Param #
random_flip_6 (RandomFlip)	(None, 32, 32, 3)	0
random_zoom_5 (RandomZoom)	(None, 32, 32, 3)	0
random_rotation_5 (RandomRotation)	(None, 32, 32, 3)	0
conv2d_21 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_21 (MaxPooling2D)	(None, 16, 16, 32)	0
dropout_28 (Dropout)	(None, 16, 16, 32)	0
conv2d_22 (Conv2D)	(None, 16, 16, 64)	18,496
max_pooling2d_22 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_29 (Dropout)	(None, 8, 8, 64)	0
conv2d_23 (Conv2D)	(None, 8, 8, 128)	73,856
max_pooling2d_23 (MaxPooling2D)	(None, 4, 4, 128)	0
dropout_30 (Dropout)	(None, 4, 4, 128)	0
flatten_8 (Flatten)	(None, 2048)	0
dense_17 (Dense)	(None, 256)	524,544
dropout_31 (Dropout)	(None, 256)	0
dense_18 (Dense)	(None, 10)	2,570

Total params: 620,362 (2.37 MB)

Trainable params: 620,362 (2.37 MB)

Non-trainable params: 0 (0.00 B)

```
In [108... early = keras.callbacks.EarlyStopping( # Callback para detener el entrenamiento tem
monitor="val_accuracy", # Monitorear la precisión en el conjunto de validación
patience=5, # Paciencia para detener el entrenamiento
restore_best_weights=True # Restaurar los mejores pesos al detener el entrenami
)

start = time.time() # Inicio del temporizador
# Entrenamiento del modelo CNN
hist_cnn_aug = cnn_aug.fit(
    x_train, y_train, # Datos de entrenamiento
    validation_split=0.2, # Usar 20% de los datos de entrenamiento para validación
    batch_size=64, # Tamaño del lote
    epochs=10, #
    callbacks=[early],
    verbose=2
)
t_train_cnn = time.time() - start # Tiempo de entrenamiento
print(f"Tiempo de entrenamiento CNN: {t_train_cnn:.2f} s")

test_loss_cnn, test_acc_cnn = cnn_aug.evaluate(x_test, y_test, verbose=0)
print(f"[CNN] Test Acc: {test_acc_cnn:.4f} | Test Loss: {test_loss_cnn:.4f}")
```

Epoch 1/10

625/625 - 36s - 58ms/step - accuracy: 0.2925 - loss: 1.8875 - val\_accuracy: 0.4099 - val\_loss: 1.6239

Epoch 2/10

625/625 - 33s - 52ms/step - accuracy: 0.4114 - loss: 1.6090 - val\_accuracy: 0.4784 - val\_loss: 1.4227

Epoch 3/10

625/625 - 33s - 53ms/step - accuracy: 0.4447 - loss: 1.5240 - val\_accuracy: 0.4923 - val\_loss: 1.3913

Epoch 4/10

625/625 - 34s - 54ms/step - accuracy: 0.4676 - loss: 1.4725 - val\_accuracy: 0.5059 - val\_loss: 1.3575

Epoch 5/10

625/625 - 35s - 57ms/step - accuracy: 0.4843 - loss: 1.4275 - val\_accuracy: 0.5449 - val\_loss: 1.2801

Epoch 6/10

625/625 - 37s - 58ms/step - accuracy: 0.4971 - loss: 1.3910 - val\_accuracy: 0.5346 - val\_loss: 1.2944

Epoch 7/10

625/625 - 35s - 56ms/step - accuracy: 0.5070 - loss: 1.3751 - val\_accuracy: 0.5292 - val\_loss: 1.2855

Epoch 8/10

625/625 - 36s - 58ms/step - accuracy: 0.5185 - loss: 1.3428 - val\_accuracy: 0.5584 - val\_loss: 1.2218

Epoch 9/10

625/625 - 36s - 58ms/step - accuracy: 0.5256 - loss: 1.3169 - val\_accuracy: 0.5528 - val\_loss: 1.2357

Epoch 10/10

625/625 - 36s - 58ms/step - accuracy: 0.5311 - loss: 1.3043 - val\_accuracy: 0.5566 - val\_loss: 1.2300

Tiempo de entrenamiento CNN: 351.83 s

[CNN] Test Acc: 0.5594 | Test Loss: 1.2195



In [109...


```

# Learning curve
plt.title("Curva de Aprendizaje - CNN with Augmented Data\n CIFAR-10")
plt.plot(np.arange(0,epochs), hist_cnn.history["loss"], label="Train")
plt.plot(np.arange(0,epochs), hist_cnn.history["val_loss"], label="Validation")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.show()

# Confusion Matrix
pred_start = time.time()
y_pred_prob = cnn_aug.predict(x_test)
print(f"Tiempo de prediccion {time.time() -pred_start}")
y_pred_class = np.argmax(y_pred_prob,axis=1)
y_test_class = np.argmax(y_test, axis=1)
print(classification_report(y_test_class, y_pred_class))
cm = confusion_matrix(y_test_class, y_pred_class)
sns.heatmap(cm, xticklabels=class_names, yticklabels=class_names)
plt.title("CNN with Augmented Data - Matriz de Confusion\n CIFAR-10")

```

Curva de Aprendizaje - CNN with Augmented Data  
CIFAR-10




313/313 ————— 3s 9ms/step

Tiempo de prediccion 3.0713276863098145

	precision	recall	f1-score	support
0	0.61	0.64	0.63	1000
1	0.54	0.82	0.65	1000
2	0.59	0.27	0.37	1000
3	0.44	0.22	0.29	1000
4	0.51	0.42	0.46	1000
5	0.59	0.41	0.48	1000
6	0.51	0.81	0.63	1000
7	0.54	0.72	0.62	1000
8	0.68	0.66	0.67	1000
9	0.57	0.62	0.60	1000
accuracy			0.56	10000
macro avg	0.56	0.56	0.54	10000
weighted avg	0.56	0.56	0.54	10000

Out[109...] Text(0.5, 1.0, 'CNN with Augmented Data - Matriz de Confusion\n CIFAR-10')



In [110...] show\_prediction\_errors(cnn\_aug, "CNN with Augmented Data", 5)

## Error Examples - CNN with Augmented Data

Real: cat

Predicted: truck



Real: ship

Predicted: automobile




Real: airplane

Predicted: ship



Real: dog

Predicted: frog



Real: ship

Predicted: deer

